

This week

- Ensemble methods
 - Bagging
 - Random forests

Today

- Bagging
 - code: `ants_bag.R`, `ants_bag.py`
 - inference algorithm (predictive performance)
 - tuning
- Parallel processing

Ensemble methods

- Train many models (ensemble)
- Average the models to predict
- Averaging reduces prediction variance

$$\text{e.g. } \text{Var}(\bar{y}) = \frac{\sigma_y^2}{n} \quad (\text{for independent } y)$$

Takeaway:

variance of the mean of y is
less than the variance of y

Bagging

- Bootstrap
 - form many new datasets by resampling from the data
 - sample with replacement
 - train model on each dataset
- Aggregate
 - average over trained models

Bagging algorithm

for many repetitions
 resample the data with replacement
 train the base model
 record prediction
final prediction = mean of predictions

Base model: can be any type of model

Bagged regression tree

Algorithm pseudocode to R

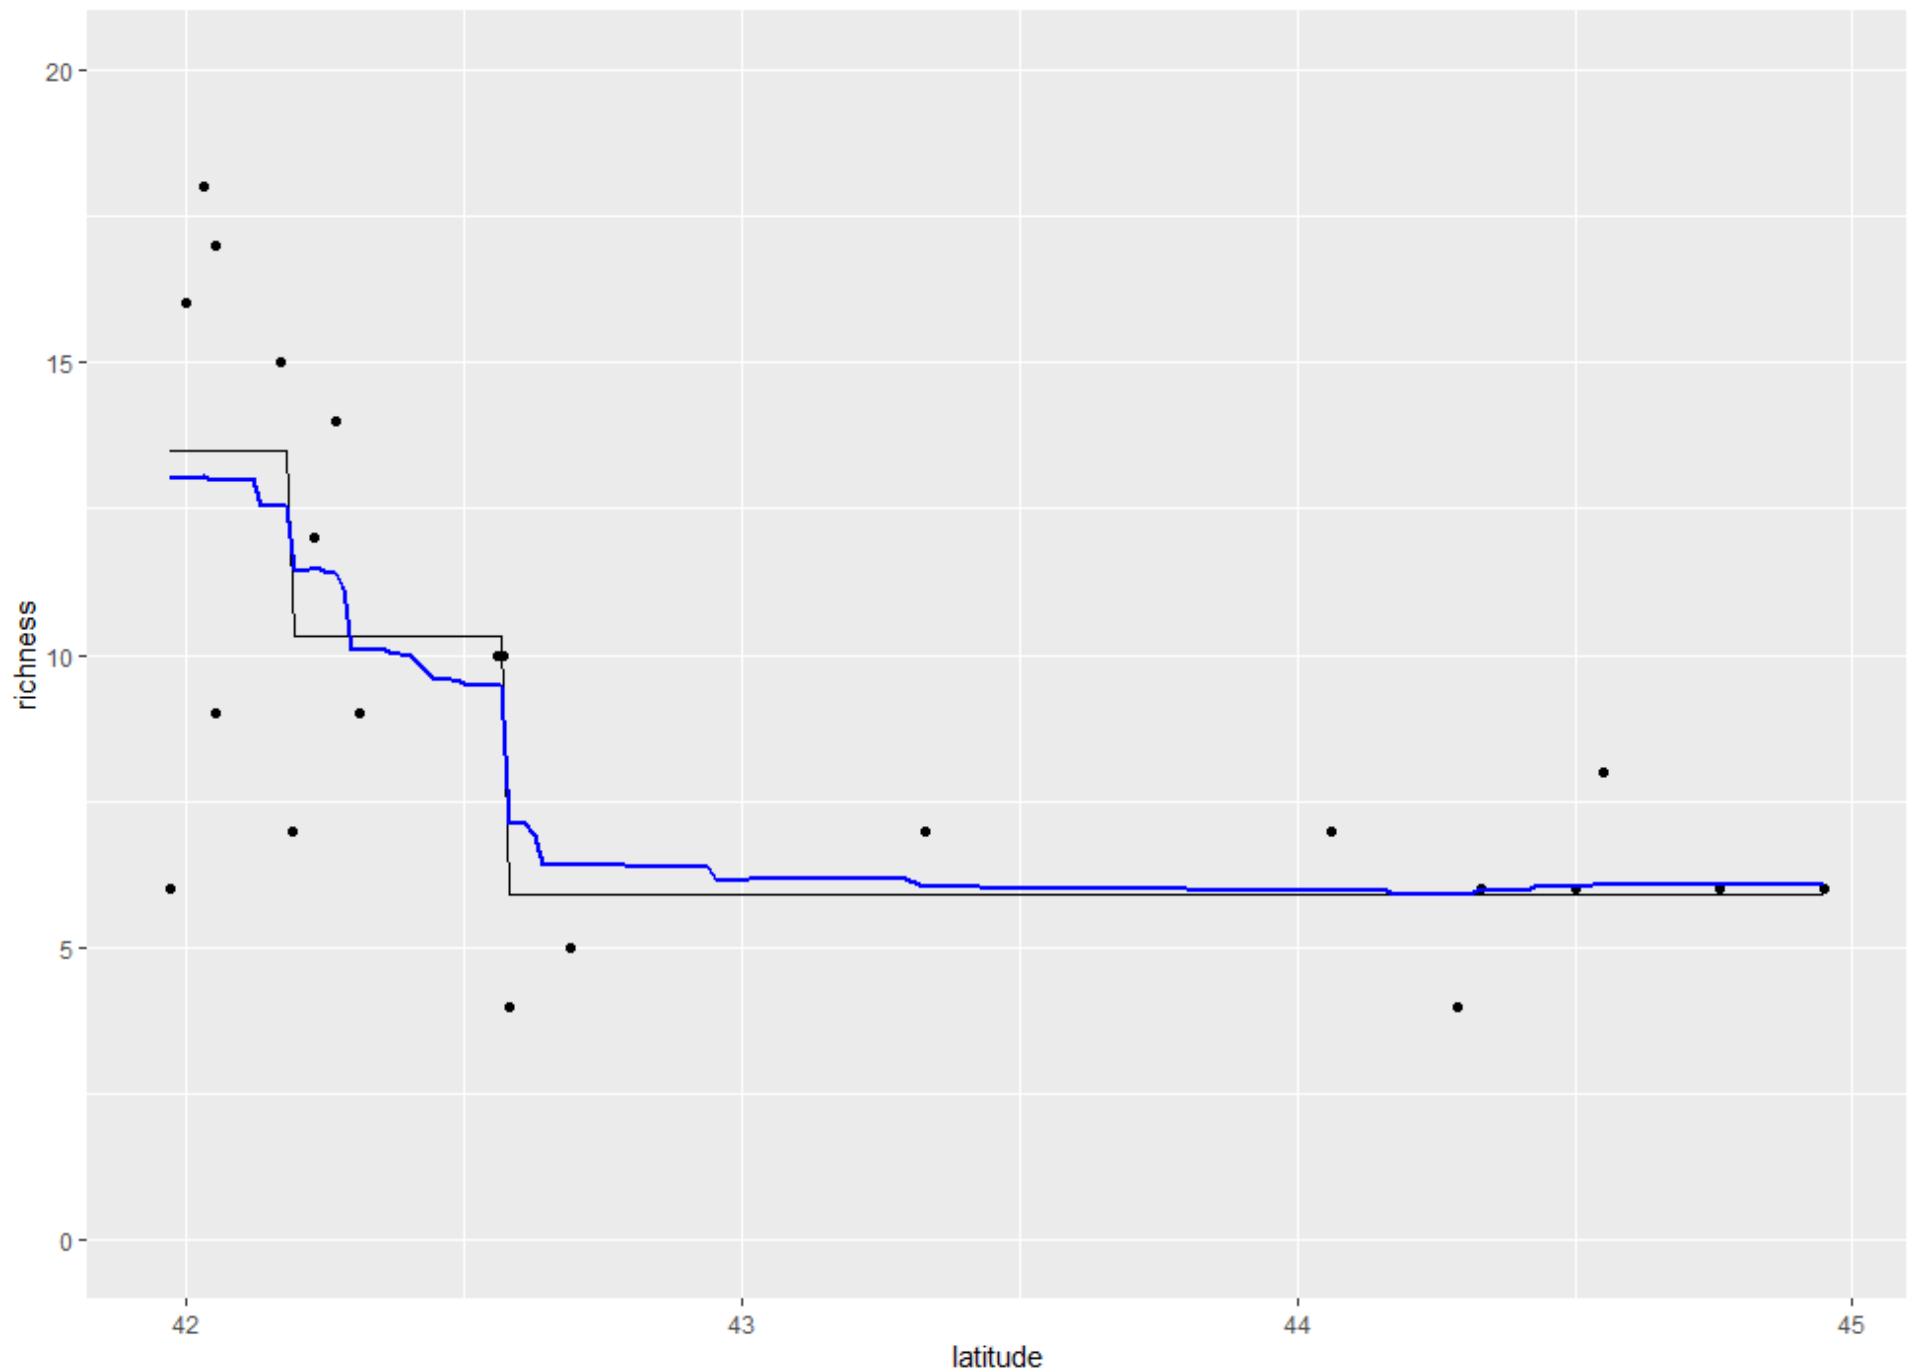
```
# Bagging algorithm
boot_reps <- 500
n <- nrow(forest_ants)
nx <- nrow(grid_data)
boot_preds <- matrix(rep(NA, nx*boot_reps), nrow=nx, ncol=boot_reps)
# for many repetitions
for ( i in 1:boot_reps ) {
  # resample the data (rows) with replacement
  boot_indices <- sample(1:n, n, replace=TRUE)
  boot_data <- forest_ants[boot_indices,]
  # train the base model
  boot_train <- tree(richness ~ latitude, data=boot_data)
  # record prediction
  boot_preds[,i] <- predict(boot_train, newdata=grid_data)
}
bagged_preds <- rowMeans(boot_preds)
```

Bagged regression tree

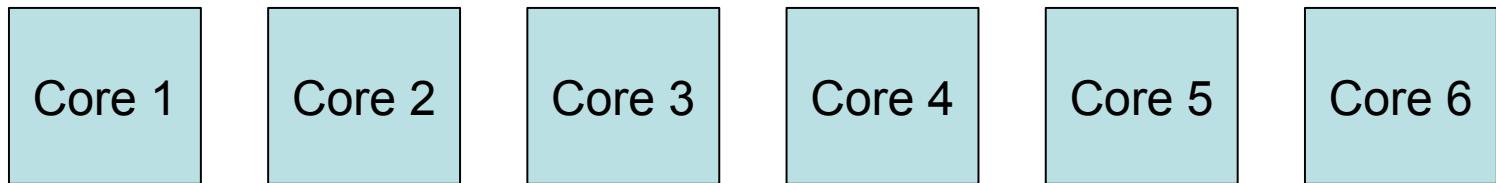
Algorithm pseudocode to Python

```
# Bagging algorithm
boot_reps = 500
dt = tree.DecisionTreeRegressor(max_depth=2) #define base model
n = len(forest_ants)
nx = len(grid_data)
boot_preds = np.full((nx, boot_reps), np.nan)
# for many repetitions
for i in range(boot_reps):
    # resample the data (rows) with replacement
    boot_indices = rng.choice(range(n), n, replace=True)
    boot_data = forest_ants.iloc[boot_indices]
    # train the base model
    boot_train = dt.fit(boot_data[["latitude"]], boot_data["richness"])
    # record prediction
    boot_preds[:,i] = boot_train.predict(grid_data)
# mean of predictions
bagged_preds = np.mean(boot_preds, axis=1)
```

Bagged regression tree (blue) vs single regression tree (black)



Parallel processing

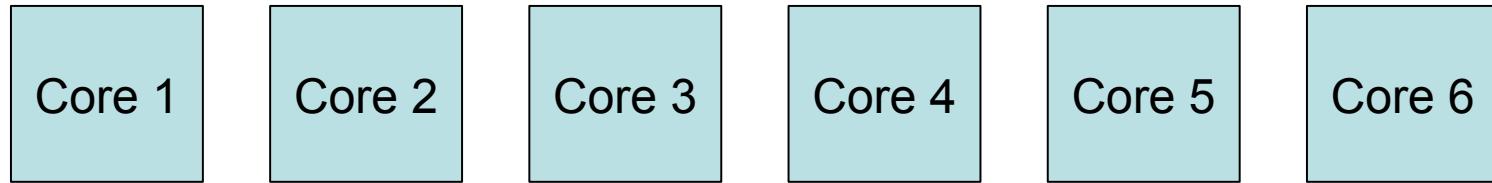


6000

sequential

```
for ( i in 1:6000 ) {  
  y[i] <- long_job()  
}
```

Parallel processing



6000
Core 1 Core 2 Core 3 Core 4 Core 5 Core 6

sequential

parallel

```
for ( i in 1:6000 ) {  
  y[i] <- long_job()  
}
```

```
y <- foreach ( i=1:6000 ) %doRNG%  
  long_job()  
}
```

Setup

```
library(doFuture)           → library(future)
library(doRNG)              library(foreach)
registerDoFuture()

availableCores()
plan(multisession, workers=8)

# Handy timing function:
system.time( some_function() )

# Saving/loading long jobs
save(myresult1, myresult2, file="/saved/myresult.Rdata")
load("/saved/myresult.Rdata")
```

Inference algorithm

- 5-fold CV, 500 splits
- Ants: mean prediction error (MSE)
12.93 +/- 0.07

Model	LOOCV	5-fold CV
Polynomial 2	12.88	13.51
Single reg tree	12.68	13.15
KNN 7	12.63	13.03
KNN 6	12.95	13.01
Bagged reg tree	13.23*	12.93
Smoothing spline 3	12.52	12.77

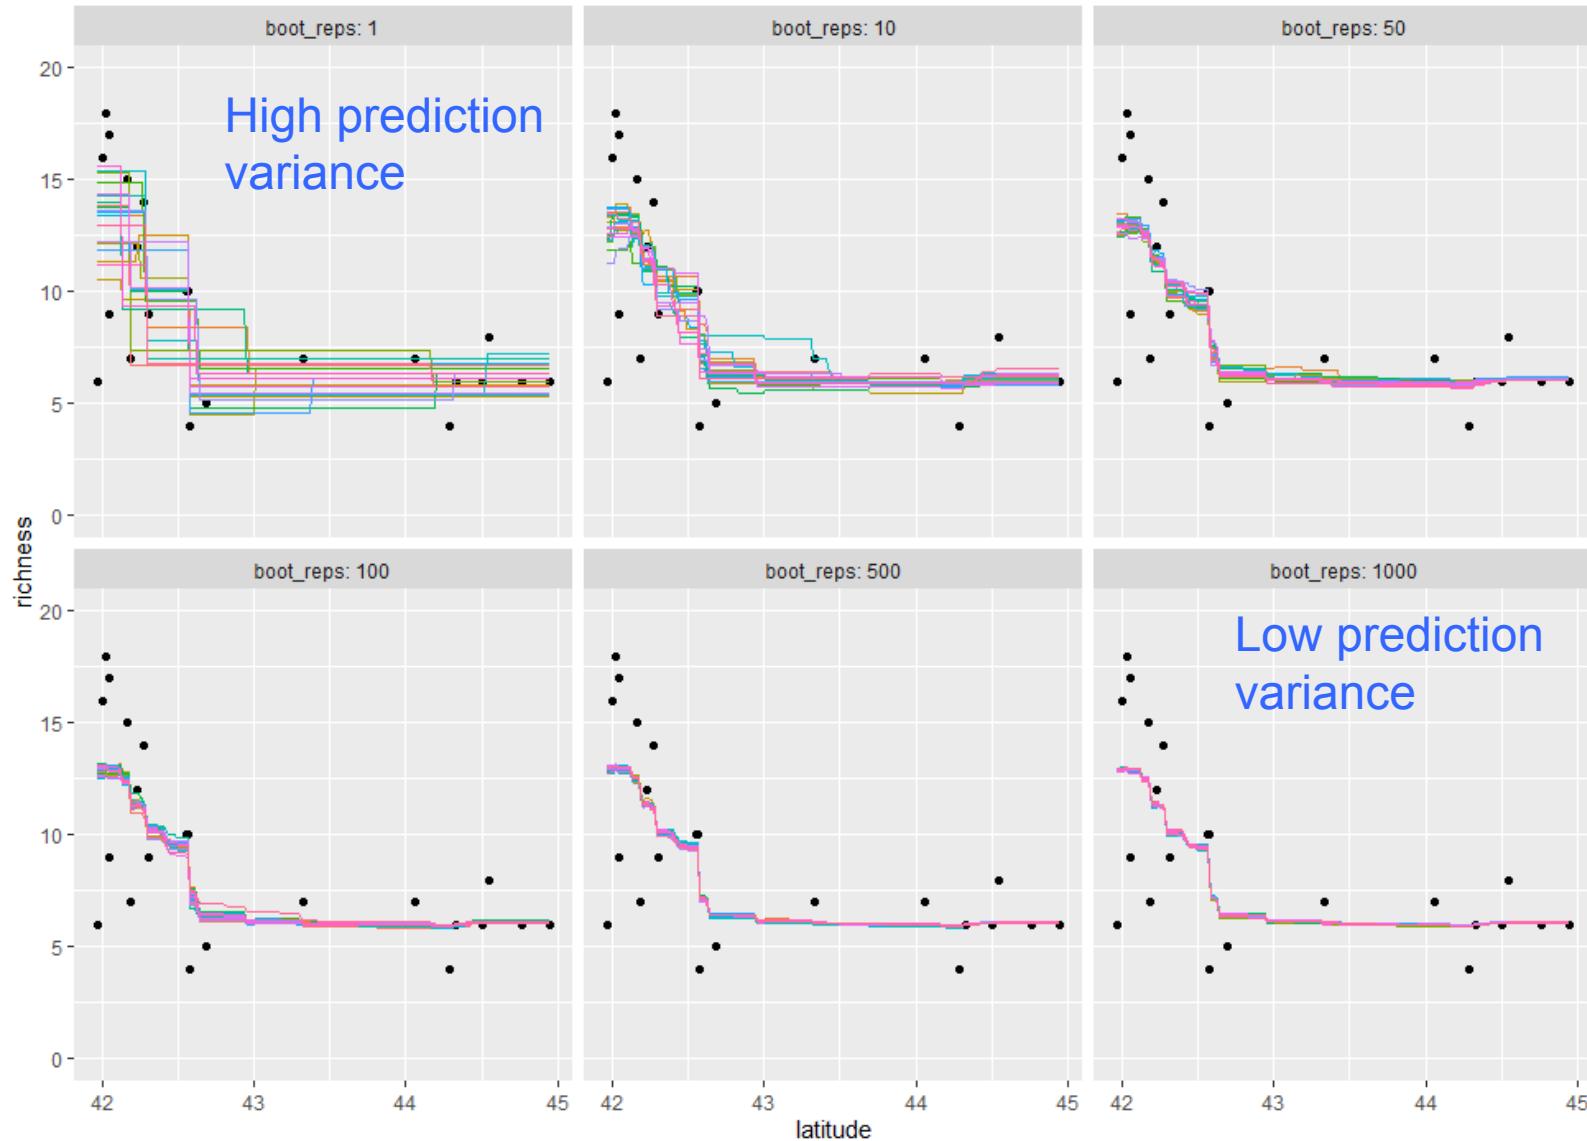
2nd best on 5-fold CV:
generally good predictive function

Worst on LOOCV:
perhaps an influential data point

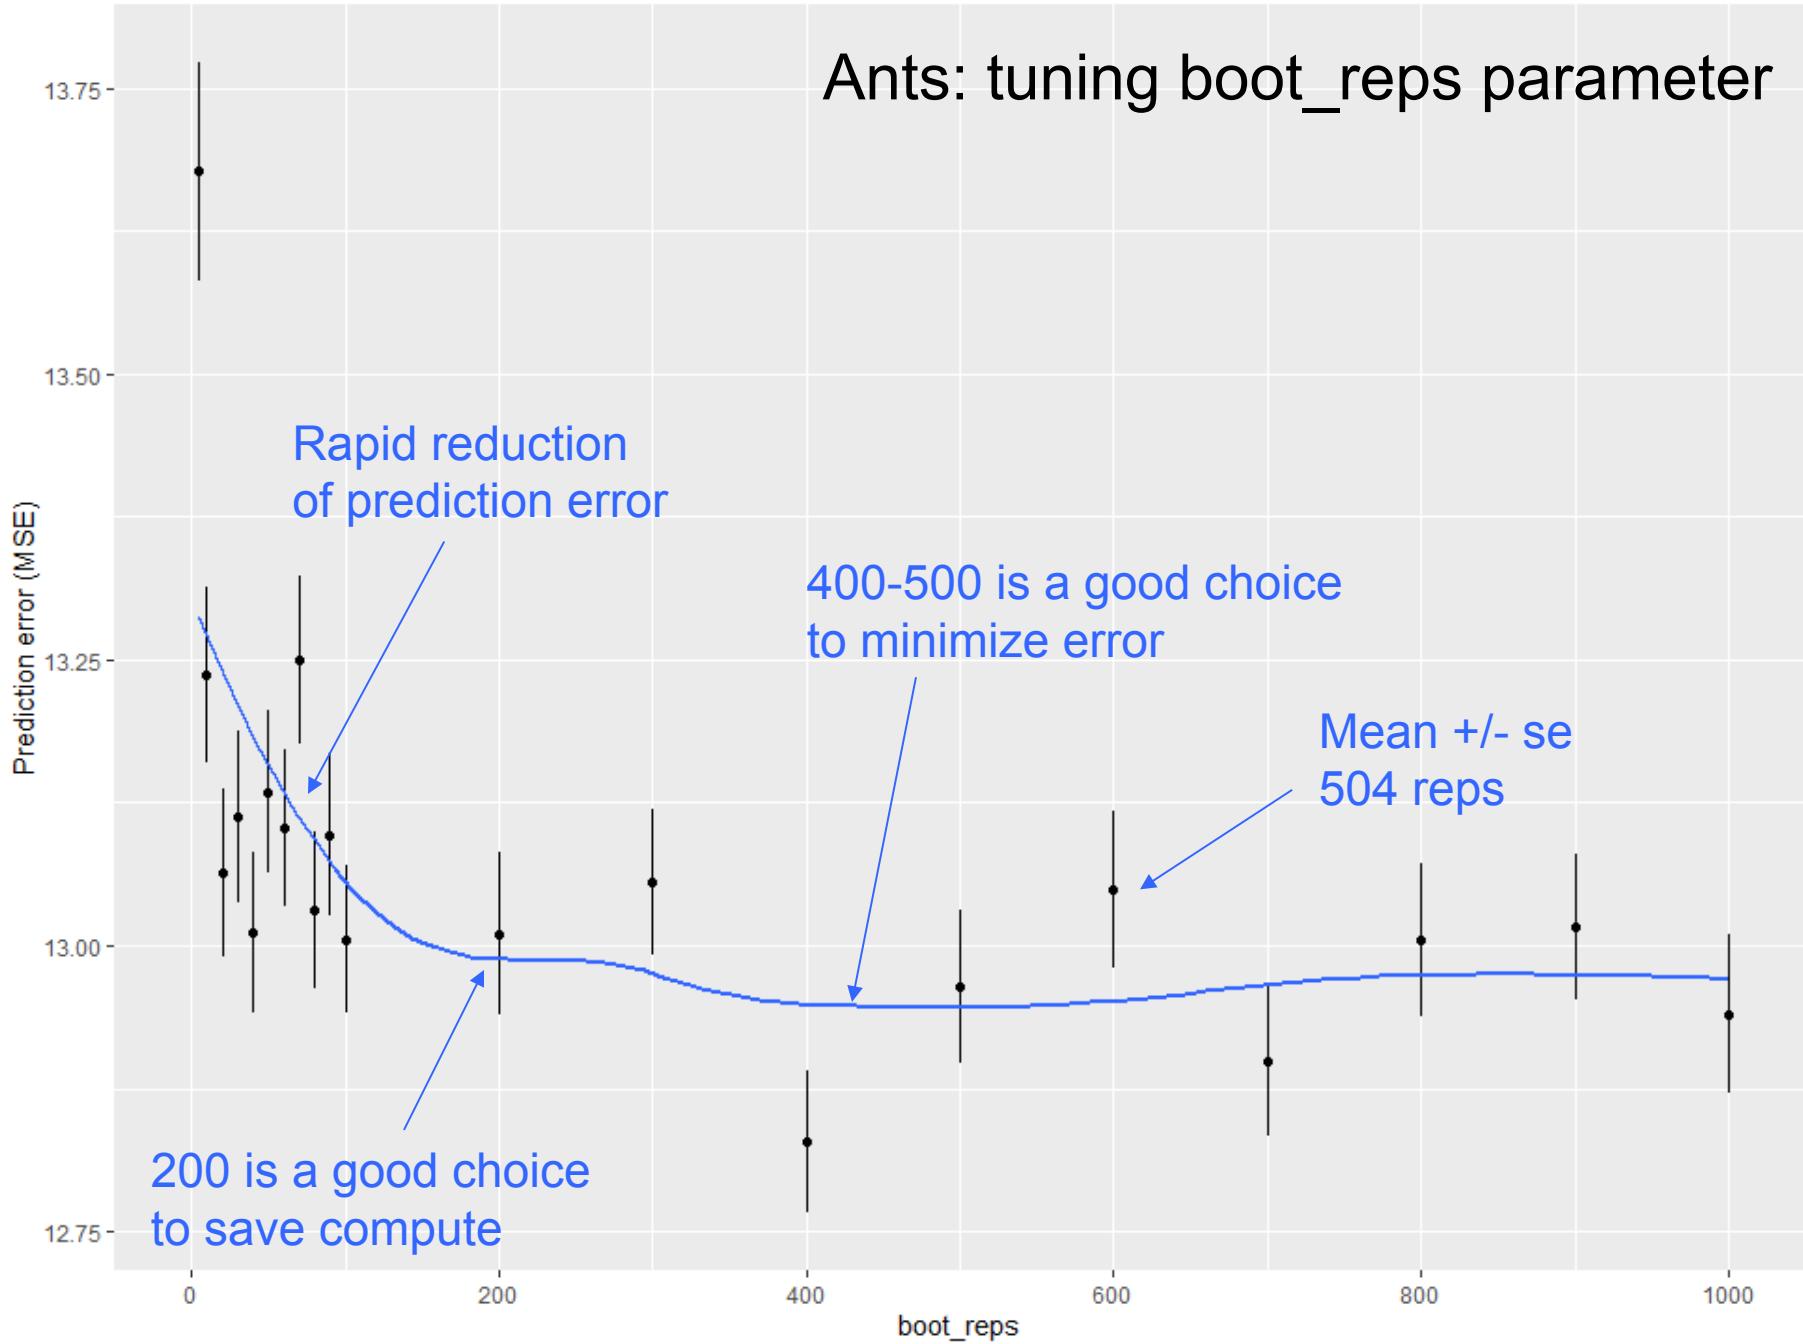
Prediction is always tenuous
with a small dataset

Bagging reduces prediction variance

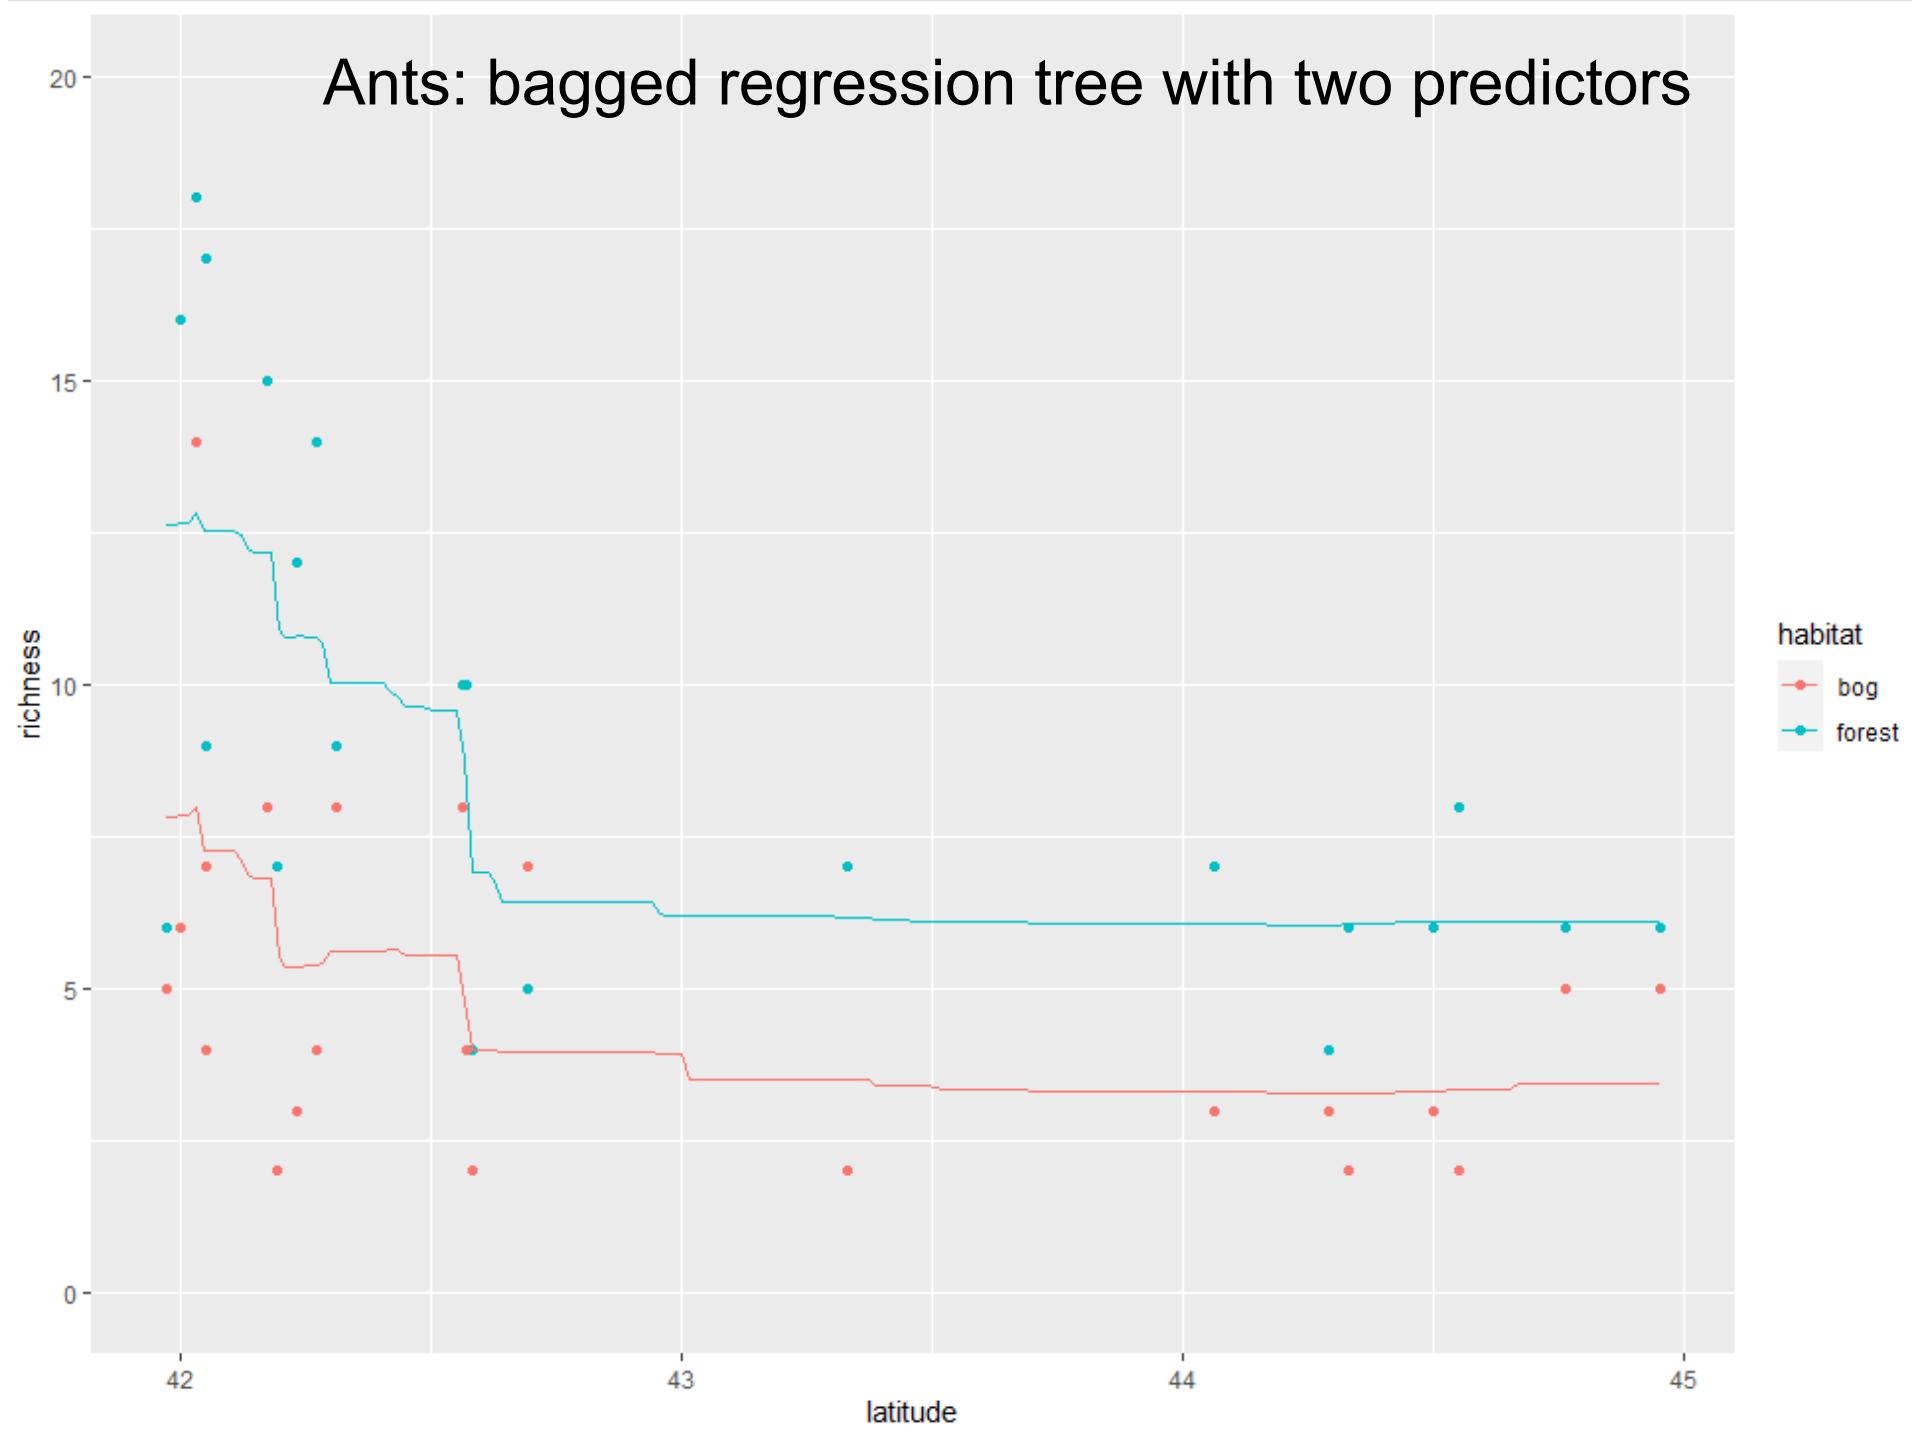
Each panel shows 20 realizations of `bagrt()`



Ants: tuning boot_reps parameter



Ants: bagged regression tree with two predictors



Bagged KNN 7 (blue) compared to single KNN 7 (black)

Bagging with KNN as base model

